

# DocRED: A Large-Scale Document-Level Relation Extraction Dataset

Yuan Yao<sup>1\*</sup>, Deming Ye<sup>1\*</sup>, Peng Li<sup>2</sup>, Xu Han<sup>1</sup>, Yankai Lin<sup>1</sup>, Zhenghao Liu<sup>1</sup>,  
Zhiyuan Liu<sup>1†</sup>, Lixin Huang<sup>1</sup>, Jie Zhou<sup>2</sup>, Maosong Sun<sup>1</sup>

<sup>1</sup>Department of Computer Science and Technology, Tsinghua University, Beijing, China  
Institute for Artificial Intelligence, Tsinghua University, Beijing, China  
State Key Lab on Intelligent Technology and Systems, Tsinghua University, Beijing, China

<sup>2</sup>Pattern Recognition Center, WeChat AI, Tencent Inc.  
{yuan-yao18, ydm18}@mails.tsinghua.edu.cn

## Abstract

**Multiple entities** in a document generally exhibit complex inter-sentence relations, and cannot be well handled by existing relation extraction (RE) methods that typically focus on extracting intra-sentence relations for single entity pairs. In order to accelerate the research on document-level RE, we introduce **DocRED**, a new dataset constructed from Wikipedia and Wikidata with three features: (1) DocRED annotates both named entities and relations, and is the **largest human-annotated dataset** for document-level RE from plain text; (2) DocRED requires reading multiple sentences in a document to **extract entities and infer their relations** by synthesizing all information of the document; (3) along with the human-annotated data, we also offer **large-scale distantly supervised data**, which enables DocRED to be adopted for both supervised and weakly supervised scenarios. In order to verify the challenges of document-level RE, we implement recent state-of-the-art methods for RE and conduct a thorough evaluation of these methods on DocRED. Empirical results show that DocRED is challenging for existing RE methods, which indicates that document-level RE remains an open problem and requires further efforts. Based on the detailed analysis on the experiments, we discuss multiple promising directions for future research. We make DocRED and the code for our baselines publicly available at <https://github.com/thunlp/DocRED>.

## 1 Introduction

The task of relation extraction (RE) is to identify relational facts between entities from plain text, which plays an important role in large-scale knowledge graph construction. Most existing RE

### Kungliga Hovkapellet

[1] *Kungliga Hovkapellet* (The *Royal Court Orchestra*) is a *Swedish* orchestra, originally part of the *Royal Court* in *Sweden*'s capital *Stockholm*. [2] The orchestra originally consisted of both musicians and singers. [3] It had only male members until 1727, when *Sophia Schröder* and *Judith Fischer* were employed as vocalists; in the 1850s, the harpist *Marie Pauline Ahman* became the first female instrumentalist. [4] From 1731, public concerts were performed at *Riddarhuset* in *Stockholm*. [5] Since 1773, when the *Royal Swedish Opera* was founded by *Gustav III* of *Sweden*, the *Kungliga Hovkapellet* has been part of the opera's company.

Subject: *Kungliga Hovkapellet*; *Royal Court Orchestra*  
Object: *Royal Swedish Opera*  
Relation: **part\_of** Supporting Evidence: 5

Subject: *Riddarhuset*  
Object: *Sweden*  
Relation: **country** Supporting Evidence: 1, 4

Figure 1: An example from DocRED. Each document in DocRED is annotated with **named entity mentions**, **coreference information**, **intra- and inter-sentence relations**, and **supporting evidence**. 2 out of the 19 relation instances annotated for this example document are presented, with named entity mentions involved in these instances colored in blue and other named entity mentions underlined for clarity. Note that mentions of the same subject (e.g., *Kungliga Hovkapellet* and *Royal Court Orchestra*) are identified as shown in the first relation instance.

work focuses on sentence-level RE, i.e., extracting relational facts from a single sentence. In recent years, various neural models have been explored to encode relational patterns of entities for sentence-level RE, and achieve state-of-the-art performance (Socher et al., 2012; Zeng et al., 2014, 2015; dos Santos et al., 2015; Xiao and Liu, 2016; Cai et al., 2016; Lin et al., 2016; Wu et al., 2017; Qin et al., 2018; Han et al., 2018a).

Despite these successful efforts, **sentence-level RE suffers from an inevitable restriction in practice: a large number of relational facts are expressed in multiple sentences**. Taking Figure 1 as an example, multiple entities are mentioned in the document and exhibit complex interactions. In

\* indicates equal contribution

† Corresponding author: Z.Liu(liuzy@tsinghua.edu.cn)

order to identify the relational fact (*Riddarhuset*, *country*, *Sweden*), one has to first identify the fact that *Riddarhuset* is located in *Stockholm* from Sentence 4, then identify the facts *Stockholm* is the capital of *Sweden* and *Sweden* is a country from Sentence 1, and finally infer from these facts that the sovereign state of *Riddarhuset* is *Sweden*. The process requires reading and reasoning over multiple sentences in a document, which is intuitively beyond the reach of sentence-level RE methods. According to the statistics on our human-annotated corpus sampled from Wikipedia documents, at least 40.7% relational facts can only be extracted from multiple sentences, which is not negligible. Swampillai and Stevenson (2010) and Verga et al. (2018) have also reported similar observations. Therefore, it is necessary to move RE forward from sentence level to document level.

The research on document-level RE requires a large-scale annotated dataset for both training and evaluation. Currently, there are only a few datasets for document-level RE. Quirk and Poon (2017) and Peng et al. (2017) build two distantly supervised datasets without human annotation, which may make the evaluation less reliable. BC5CDR (Li et al., 2016) is a human-annotated document-level RE dataset consisting of 1,500 PubMed documents, which is in the specific domain of biomedicine considering only the “chemical-induced disease” relation, making it unsuitable for developing general-purpose methods for document-level RE. Levy et al. (2017) extract relational facts from documents by answering questions using reading comprehension methods, where the questions are converted from entity-relation pairs. As the dataset proposed in this work is tailored to the specific approach, it is also unsuitable for other potential approaches for document-level RE. In summary, existing datasets for document-level RE either only have a small number of manually-annotated relations and entities, or exhibit noisy annotations from distant supervision, or serve specific domains or approaches. In order to accelerate the research on document-level RE, we urgently need a large-scale, manually-annotated, and general-purpose document-level RE dataset.

In this paper, we present DocRED, a large-scale human-annotated document-level RE dataset constructed from Wikipedia and Wikidata (Erxleben et al., 2014; Vrandečić and Krötzsch, 2014). Do-

docRED is constructed with the following **three features**: (1) DocRED contains 132,375 entities and 56,354 relational facts annotated on 5,053 Wikipedia documents, making it the largest human-annotated document-level RE dataset. (2) As at least 40.7% of the relational facts in DocRED can only be extracted from multiple sentences, DocRED requires reading multiple sentences in a document to recognize entities and inferring their relations by synthesizing all information of the document. This distinguishes DocRED from those sentence-level RE datasets. (3) We also provide large-scale distantly supervised data to support weakly supervised RE research.

To assess the challenges of DocRED, we implement recent state-of-the-art RE methods and conduct thorough experiments on DocRED under various settings. Experimental results show that the performance of existing methods declines significantly on DocRED, indicating the task document-level RE is more challenging than sentence-level RE and remains an open problem. Furthermore, detailed analysis on the results also reveals multiple promising directions worth pursuing.

## 2 Data Collection

Our ultimate goal is to construct a dataset for document-level RE from plain text, which requires necessary information including **named entity mentions**, **entity coreferences**, and **relations of all entity pairs** in the document. To facilitate more RE settings, we also provide supporting **evidence information for relation instances**. In the following sections, we first introduce the collection process of the human-annotated data, and then describe the process of creating the large-scale distantly supervised data.

### 2.1 Human-Annotated Data Collection

Our human-annotated data is collected in four stages: (1) Generating distantly supervised annotation for Wikipedia documents. (2) Annotating all named entity mentions in the documents and coreference information. (3) Linking named entity mentions to Wikidata items. (4) Labeling relations and corresponding supporting evidence.

Following ACE annotation process (Dodington et al., 2004), both Stage 2 and 4 require three iterative passes over the data: (1) **Generating named entity** using named entity recognition (NER) models, or **relation recommendations** us-

ing distant supervision and RE models. (2) Manually correcting and supplementing recommendations. (3) Reviewing and further modifying the annotation results from the second pass for better accuracy and consistency. To ensure the annotators are well trained, a principled training procedure is adopted and the annotators are required to pass test tasks before annotating the dataset. And only carefully selected experienced annotators are qualified for the third pass annotation.

To provide a strong alignment between text and KBs, our dataset is constructed from the complete English Wikipedia document collection and Wikidata<sup>1</sup>, which is a large-scale KB tightly integrated with Wikipedia. We use the introductory sections from Wikipedia documents as the corpus, as they are usually high-quality and contain most of the key information.

**Stage 1: Distantly Supervised Annotation Generation.** To select documents for human annotation, we align Wikipedia documents with Wikidata under the distant supervision assumption (Mintz et al., 2009). Specifically, we first perform named entity recognition using spaCy<sup>2</sup>. Then these named entity mentions are linked to Wikidata items, where named entity mentions with identical KB IDs are merged. Finally, relations between each merged named entity pair in the document are labeled by querying Wikidata. Documents containing fewer than 128 words are discarded. To encourage reasoning, we further discard documents containing fewer than 4 entities or fewer than 4 relation instances, resulting in 107,050 documents with distantly supervised labels, where we randomly select 5,053 documents and the most frequent 96 relations for human annotation.

**Stage 2: Named Entity and Coreference Annotation.** Extracting relations from document requires first recognizing named entity mentions and identifying mentions referring to the same entities within the document. To provide high-quality named entity mentions and coreference information, we ask human annotators first to review, correct and supplement the named entity mention recommendations generated in Stage 1, and then merge those different mentions referring to the same entities, which provides extra coreference information. The resulting intermediate corpus con-

tains a variety of named entity types including person, location, organization, time, number and names of miscellaneous entities that do not belong to the aforementioned types.

**Stage 3: Entity Linking.** In this stage, we link each named entity mention to multiple Wikidata items to provide relation recommendations from distant supervision for the next stage. To be specific, each named entity mention is associated with a Wikidata item candidate set<sup>3</sup> consisting of all Wikidata items whose names or aliases literally match it. We further extend the candidate set using Wikidata items hyperlinked to the named entity mention by the document authors, and recommendations from an entity linking toolkit TagMe (Ferragina and Scaiella, 2010). Specially, numbers and time are semantically matched.

**Stage 4: Relation and Supporting Evidence Collection.** The annotation of relation and supporting evidence is based on the named entity mentions and coreference information in Stage 2, and faces two main challenges. The first challenge comes from the large number of potential entity pairs in the document. On the one hand, given the quadratic number of potential entity pairs with regard to entity number (19.5 entities on average) in a document, exhaustively labeling relations between each entity pair would lead to intensive workload. On the other hand, most entity pairs in a document do not contain relations. The second challenge lies in the large number of fine-grained relation types in our dataset. Thus it is not feasible for annotators to label relations from scratch.

We address the problem by providing human annotators with recommendations from RE models, and distant supervision based on entity linking (Stage 3). On average, we recommend 19.9 relation instances per document from entity linking, and 7.8 from RE models for supplement. We ask the annotators to review the recommendations, remove the incorrect relation instances and supplement omitted ones. We also ask the annotators to further select all sentences that support the reserved relation instances as supporting evidence. Relations reserved must be reflected in the document, without relying on external world knowledge. Finally 57.2% relation instances from entity linking and 48.2% from RE models are reserved.

<sup>1</sup>We use the 2018-5-24 dump of English Wikipedia and 2018-3-20 dump of Wikidata.

<sup>2</sup><https://spacy.io>

<sup>3</sup>To avoid losing relation recommendations due to prediction errors in entity linking, we include multiple linking results from different approaches in the candidate set.

Dataset	# Doc.	# Word	# Sent.	# Ent.	# Rel.	# Inst.	# Fact
SemEval-2010 Task 8	-	205k	10,717	21,434	9	8,853	8,383
ACE 2003-2004	-	297k	12,783	46,108	24	16,771	16,536
TACRED	-	1,823k	53,791	152,527	41	21,773	5,976
FewRel	-	1,397k	56,109	72,124	100	70,000	55,803
BC5CDR	1,500	282k	11,089	29,271	1	3,116	2,434
DocRED (Human-annotated)	5,053	1,002k	40,276	132,375	96	63,427	56,354
DocRED (Distantly Supervised)	101,873	21,368k	828,115	2,558,350	96	1,508,320	881,298

Table 1: Statistics of RE datasets (Doc.: document, Sent.: sentence, Ent.: entity, Rel.: relation type, Inst.: relation instance, Fact: relational fact). The first four are sentence-level RE datasets.

## 2.2 Distantly Supervised Data Collection

In addition to the human-annotated data, we also collect large-scale distantly supervised data to promote weakly supervised RE scenarios. We remove the 5,053 human-annotated documents from the 106,926 documents, and use the rest 101,873 documents as the corpus of distantly supervised data. To ensure that the distantly supervised data and human-annotated data share the same entity distribution, named entity mentions are re-identified using Bidirectional Encoder Representations from Transformers (BERT) (Devlin et al., 2019) that is fine-tuned on the human-annotated data collected in Sec. 2.1 and achieves 90.5% F1 score. We link each named entity mention to one Wikidata item by a heuristic-based method, which jointly considers the frequency of a target Wikidata item and its relevance to the current document. Then we merge the named entity mentions with identical KB IDs. Finally, relations between each merged entity pair are labeled via distant supervision.

## 3 Data Analysis

In this section, we analyze various aspects of DocRED to provide a deeper understanding of the dataset and the task of document-level RE.

**Data Size.** Table 1 shows statistics of DocRED and some representative RE datasets, including sentence-level RE datasets SemEval-2010 Task 8 (Hendrickx et al., 2010), ACE 2003-2004 (Dodgington et al., 2004), TACRED (Zhang et al., 2017), FewRel (Han et al., 2018b) and document-level RE dataset BC5CDR (Li et al., 2016). We find that DocRED is larger than existing datasets in many aspects, including the number of documents, words, sentences, entities, especially in aspects of relation types, relation instances and relational facts. We hope the large-scale DocRED dataset could drive relation extraction from sen-

tence level forward to document level.

**Named Entity Types.** DocRED covers a variety of **entity types**, including person (18.5%), location (30.9%), organization (14.4%), time (15.8%) and number (5.1%). It also covers a diverse set of miscellaneous entity names (15.2%) not belonging to the aforementioned types, such as events, artistic works and laws. Each entity is annotated with 1.34 mentions on average.

**Relation Types.** Our dataset includes 96 frequent relation types from Wikidata. A notable property of our dataset is that the **relation types** cover a broad range of categories, including relations relevant to science (33.3%), art (11.5%), time (8.3%), personal life (4.2%), etc., which means the relational facts are not constrained in any specific domain. In addition, the relation types are organized in a well-defined hierarchy and taxonomy, which could provide rich information for document-level RE systems.

**Reasoning Types.** We randomly sampled 300 documents from dev and test set, which contain 3,820 relation instances, and manually analyze the reasoning types required to extract these relations. Table 2 shows statistics of major reasoning types in our dataset. From the statistics on reasoning types, we have the following observations: (1) Most of the relation instances (61.1%) require reasoning to be identified, and only 38.9% relation instances can be extracted via simple pattern recognition, which indicates that reasoning is essential for document-level RE. (2) In relation instances with reasoning, a majority (26.6%) require logical reasoning, where the relations between two entities in question are indirectly established by a bridge entity. **Logical reasoning requires RE systems to be capable of modeling interactions between multiple entities.** (3) A notable number of relation instances (17.6%) need **coreference reasoning**, where coreference resolution must be performed first to identify target entities in a rich con-



Reasoning Types	%	Examples
Pattern recognition	38.9	[1] <i>Me Musical Nephews</i> is a 1942 one-reel animated cartoon directed by Seymour Kneitel and animated by Tom Johnson and George Germanetti. [2] Jack Mercer and Jack Ward wrote the script. ... <b>Relation:</b> <i>publication_date</i> <b>Supporting Evidence:</b> 1
Logical reasoning	26.6	[1] “Nisei” is the ninth episode of the third season of the American science fiction television series <i>The X-Files</i> . ... [3] It was directed by David Nutter, and written by Chris Carter, Frank Spotnitz and Howard Gordon. ... [8] The show centers on FBI special agents <i>Fox Mulder</i> (David Duchovny) and Dana Scully (Gillian Anderson) who work on cases linked to the paranormal, called X-Files. ... <b>Relation:</b> <i>creator</i> <b>Supporting Evidence:</b> 1, 3, 8
Coreference reasoning	17.6	[1] <i>Dwight Tillery</i> is an American politician of the Democratic Party who is active in local politics of Cincinnati, Ohio. ... [3] He also holds a law degree from the <i>University of Michigan Law School</i> . [4] <i>Tillery</i> served as mayor of Cincinnati from 1991 to 1993. <b>Relation:</b> <i>educated_at</i> <b>Supporting Evidence:</b> 1, 3
Common-sense reasoning	16.6	[1] <i>William Busac</i> (1020-1076), son of William I, Count of Eu, and his wife Lesceline. ... [4] <i>William</i> appealed to King Henry I of France, who gave him in marriage <i>Adelaide</i> , the heiress of the county of Soissons. [5] <i>Adelaide</i> was daughter of Renaud I, Count of Soissons, and Grand Master of the Hotel de France. ... [7] <i>William</i> and <i>Adelaide</i> had four children: ... <b>Relation:</b> <i>spouse</i> <b>Supporting Evidence:</b> 4, 7

Table 2: Types of reasoning required for document-level RE on DocRED. The rest 0.3% requires other types of reasoning, such as temporal reasoning. The *head*, *tail* and *relation* are colored accordingly.

text. (4) A similar proportion of relation instances (16.6%) has to be identified based on *common-sense reasoning*, where readers need to combine relational facts from the document with common-sense to complete the relation identification. In summary, DocRED requires rich reasoning skills for synthesizing all information of the document.

**Inter-Sentence Relation Instances.** We find that each relation instance is associated with 1.6 supporting sentences on average, where 46.4% relation instances are associated with more than one supporting sentence. Moreover, detailed analysis reveals that 40.7% relational facts can only be extracted from multiple sentences, indicating that DocRED is a good benchmark for document-level RE. We can also conclude that the abilities of reading, synthesizing and reasoning over multiple sentence are essential for document-level RE.

## 4 Benchmark Settings

We design two benchmark settings for *supervised and weakly supervised* scenarios respectively. For both settings, RE systems are evaluated on the high-quality human-annotated dataset, which provides more reliable evaluation results for document-level RE systems. The statistics of data used for the two settings are shown in Table 3.

**Supervised Setting.** In this setting, only human-annotated data is used, which are randomly split

	Setting	# Doc.	# Rel.	# Inst.	# Fact
Train	W	101,873	96	1,508,320	881,298
	S	3,053	96	38,269	34,715
Dev	S,W	1,000	96	12,332	11,790
Test	S,W	1,000	96	12,842	12,101

Table 3: Statistics of data used for the two benchmark settings (Sec. 4): supervised setting (S) and weakly supervised setting (W).

into training, development and test sets. *The supervised setting brings up two challenges for document-level RE systems as follows:*

The first challenge comes from the rich reasoning skills required for performing document-level RE. As shown in Sec. 3, about 61.1% relation instances depend on complex reasoning skills other than pattern recognition to be extracted, which requires RE systems to step beyond recognizing simple patterns in a single sentence, and reason over global and complex information in a document.

The second challenge lies in the high computational cost of modeling long documents and the massive amount of potential entity pairs in a document, which is quadratic with regard to entity number (19.5 entities on average) in a document. As a result, RE systems that model context information with algorithms of quadratic or even higher computational complexity such as (Sorokin

and Gurevych, 2017; Christopoulou et al., 2018) are not efficient enough for document-level RE. Thus the efficiency of context-aware RE systems needs to be further improved to be applicable in document-level RE.

**Weakly Supervised Setting.** This setting is identical to the supervised setting except that the training set is replaced with the distantly supervised data (Sec. 2.2). In addition to the aforementioned two challenges, the inevitable wrong labeling problem accompanied with distantly supervised data is a major challenge for RE models under weakly supervised setting. Many efforts have been devoted to alleviating the wrong labeling problem in sentence-level RE (Riedel et al., 2010; Hoffmann et al., 2011; Surdeanu et al., 2012; Lin et al., 2016). However, noise in document-level distantly supervised data is significantly more than its counterpart in sentence-level. For example, for the recommended relation instances whose head and tail entities co-occur in the same sentence (i.e. intra-sentence relation instance) in Stage 4 of human-annotated data collection (Sec. 2.1), 41.4% are labeled as incorrect, while 61.8% inter-sentence relation instances are labeled as incorrect, indicating the wrong labeling problem is more challenging for weakly supervised document-level RE. Therefore, we believe offering distantly supervised data in DocRED will accelerate the development of distantly supervised methods for document-level RE. Moreover, it is also possible to jointly leverage distantly supervised data and human-annotated data to further improve the performance of RE systems.

## 5 Experiments

To assess the challenges of DocRED, we conduct comprehensive experiments to evaluate state-of-the-art RE systems on the dataset. Specifically, we conduct experiments under both supervised and weakly supervised benchmark settings. We also assess human performance and analyze the performance for different supporting evidence types. In addition, we conduct ablation study to investigate the contribution of different features. Through detailed analysis, we discuss several future directions for document-level RE.

**Models.** We adapt four state-of-the-art RE models to document-level RE scenario, including a CNN (Zeng et al., 2014) based model, an LSTM (Hochreiter and Schmidhuber, 1997) based

model, a bidirectional LSTM (BiLSTM) (Cai et al., 2016) based model and the Context-Aware model (Sorokin and Gurevych, 2017) originally designed for leveraging contextual relations to improve intra-sentence RE. The first three models differ only at the encoder used for encoding the document and will be explained in detail in the rest of this section. We refer the readers to the original paper for the details of the Context-Aware model for space limitation.

The CNN/LSTM/BiLSTM based models first encode a document  $\mathcal{D} = \{w_i\}_{i=1}^n$  consisting of  $n$  words into a hidden state vector sequence  $\{h_i\}_{i=1}^n$  with CNN/LSTM/BiLSTM as encoder, then compute the representations for entities, and finally predict relations for each entity pair.

For each word, the features fed to the encoder is the concatenation of its GloVe word embedding (Pennington et al., 2014), entity type embedding and coreference embedding. The entity type embedding is obtained by mapping the entity type (e.g., PER, LOC, ORG) assigned to the word into a vector using an embedding matrix. The entity type is assigned by human for the human-annotated data, and by a fine-tuned BERT model for the distantly supervised data. Named entity mentions corresponding to the same entity are assigned with the same entity id, which is determined by the order of its first appearance in the document. And the entity ids are mapped into vectors as the coreference embeddings.

For each named entity mention  $m_k$  ranging from the  $s$ -th word to the  $t$ -th word, we define its representation as  $\mathbf{m}_k = \frac{1}{t-s+1} \sum_{j=s}^t \mathbf{h}_j$ . And the representation of an entity  $e_i$  with  $K$  mentions is computed as the average of the representations of these mentions:  $\mathbf{e}_i = \frac{1}{K} \sum_k \mathbf{m}_k$ .

We treat relation prediction as a multi-label classification problem. Specially, for each entity pair  $(e_i, e_j)$ , we first concatenate the entity representations with relative distance embeddings, and then use a bilinear function to compute the probability for each relation type:

$$\hat{\mathbf{e}}_i = [\mathbf{e}_i; \mathbf{E}(d_{ij})], \hat{\mathbf{e}}_j = [\mathbf{e}_j; \mathbf{E}(d_{ji})] \quad (1)$$

$$P(r|e_i, e_j) = \text{sigmoid}(\hat{\mathbf{e}}_i^T \mathbf{W}_r \hat{\mathbf{e}}_j + b_r) \quad (2)$$

where  $[\cdot; \cdot]$  denotes concatenation,  $d_{ij}$  and  $d_{ji}$  are the relative distances of the first mentions of the two entities in the document,  $\mathbf{E}$  is an embedding matrix,  $r$  is a relation type, and  $\mathbf{W}_r, b_r$  are relation type dependent trainable parameters.

Model	Dev				Test			
	Ign F1	Ign AUC	F1	AUC	Ign F1	Ign AUC	F1	AUC
Supervised Setting								
CNN	41.58	36.85	43.45	39.39	40.33	36.24	42.26	38.91
LSTM	48.44	46.62	50.68	49.48	47.71	46.27	50.07	49.25
BiLSTM	48.87	<b>47.61</b>	50.94	<b>50.26</b>	<b>48.78</b>	<b>47.61</b>	<b>51.06</b>	<b>50.43</b>
Context-Aware	<b>48.94</b>	47.22	<b>51.09</b>	50.17	48.40	46.54	50.70	49.64
Weakly Supervised Setting								
CNN	33.24	23.17	42.76	37.99	32.33	21.83	42.00	36.84
LSTM	39.37	22.39	49.92	42.79	38.27	21.74	48.88	41.35
BiLSTM	<b>41.44</b>	<b>23.21</b>	<b>51.72</b>	<b>44.44</b>	39.15	<b>22.14</b>	49.80	<b>42.87</b>
Context-Aware	40.47	22.56	51.39	43.00	<b>39.16</b>	21.58	<b>50.12</b>	41.51

Table 4: Performance of different RE models on DocRED (%).

**Evaluation Metrics.** Two widely used **metrics F1 and AUC** are used in our experiments. However, some relational facts present in both the training and dev/test sets, thus a model may memorize their relations during training and achieve a better performance on the dev/test set in an undesirable way, introducing **evaluation bias**. However, the overlap in relational facts between the training and dev/test sets is inevitable, since many common relational facts are likely to be shared in different documents. Therefore, we also report the F1 and AUC scores excluding those relational facts shared by the training and dev/test sets, denoted as Ign F1 and Ign AUC, respectively.

**Model Performance.** Table 4 shows the experimental results under the supervised and weakly supervised settings, from which **we have the following observations**: (1) Models trained with human-annotated data generally outperform their counterparts trained on distantly supervised data. This is because although large-scale distantly supervised data can be easily obtained via distant supervision, the wrong-labeling problem may harm the performance of RE systems, which makes weakly supervised setting a more difficult scenario. (2) An interesting exception is that LSTM, BiLSTM and Context-Aware trained on distantly supervised data achieve comparable F1 scores as those trained on human-annotated data but significantly lower scores on the other metrics, indicating that the overlap entity pairs between training and dev/test sets indeed cause evaluation biases. Therefore, reporting Ign F1 and Ign AUC is necessary. (3) Models leveraging rich contextual information generally achieve better performances. LSTM and BiLSTM outperform CNN, indicating the effectiveness of modeling long-dependency semantics in

Method	RE			RE+Sup		
	P	R	F1	P	R	F1
Model	55.6	52.6	54.1	46.4	43.1	44.7
Human	<b>89.7</b>	<b>86.3</b>	<b>88.0</b>	<b>71.2</b>	<b>75.8</b>	<b>73.4</b>

Table 5: Human performance (%).

document-level RE. Context-Aware achieves competitive performance, however, it cannot significantly outperform other neural models. It indicates that it is beneficial to consider the association of multiple relations in document-level RE, whereas the current models are not capable of utilizing inter-relation information well.

**Human Performance.** To assess human performance on the task of document-level RE on DocRED, we randomly sample 100 documents from the test set and ask additional crowd-workers to identify relation instances and supporting evidence. Relation instances identified in the same way as Sec. 2.1 are recommended to the crowd-workers to assist them. The original annotation results collected in Sec. 2.1 are used as ground truth. We also propose another subtask of jointly identifying relation instances and supporting evidence, and also design a pipeline model. Table 5 shows the performance of RE model and human. Humans achieve competitive results on both the document-level RE task (RE) and the jointly identifying relation and supporting evidence task (RE+Sup), indicating both the ceiling performance on DocRED and the inter-annotator agreement are relatively high. In addition, the overall performance of RE models is significantly lower than human performance, which indicates document-level RE is a challenging task, and suggests ample opportunity for improvement.

### Performance v.s. Supporting Evidence Types.

Document-level RE requires synthesizing information from multiple supporting sentences. To investigate the difficulty of synthesizing information from different types of supporting evidence, we divide the 12,332 relation instances in development set into three disjoint subsets: (1) 6,115 relation instances with only one supporting sentence (denoted as *single*); (2) 1,062 relation instances with multiple supporting sentences and the entity pair co-occur in at least one supporting sentence (denoted as *mix*); (3) 4,668 relation instances with multiple supporting sentences and the entity pair do not co-occur in any supporting sentence, which means they can only be extracted from multiple supporting sentences (denoted as *multiple*). It should be noted that when a model predicts a wrong relation, we do not know which sentences have been used as supporting evidence, thus the predicted relation instance cannot be classified into the aforementioned subsets and computing precision is infeasible. Therefore, we only report recall of the RE model for each subset, which is 51.1% for *single*, 49.4% for *mix*, and 46.6% for *multiple*. This indicates that while multiple supporting sentences in *mix* may provide complementary information, it is challenging to effectively synthesize the rich global information. Moreover, the poor performance on *multiple* suggests that RE models still struggle in extracting inter-sentence relations.

**Feature Ablations.** We conduct feature ablation studies on the BiLSTM model to investigate the contribution of different features in document-level RE, including entity types, coreference information, and the relative distance between entities (Eq. 1). Table 6 shows that the aforementioned features all have a contribution to the performance. Specifically, entity types contribute most due to their constraint on viable relation types. Coreference information and the relative distance between entities are also important for synthesizing information from multiple named entity mentions. This indicates that it is important for RE systems to leverage rich information at document level.

**Supporting Evidence Prediction.** We propose a new task to predict the supporting evidence for relation instances. On the one hand, jointly predicting the evidence provides better explainability. On the other hand, identifying supporting evidence and reasoning relational facts from text are nat-

Setting	Ign F1	Ign AUC	F1	AUC
BiLSTM	<b>48.87</b>	<b>47.61</b>	<b>50.94</b>	<b>50.26</b>
- entity type	46.81	44.46	48.70	47.29
- coreference	47.22	44.72	49.37	47.49
- distance	47.94	45.57	50.19	48.43
- all features	44.08	39.94	46.52	43.18

Table 6: Feature ablations on dev set (%).

urally dual tasks with potential mutual enhancement. We design two supporting evidence prediction methods: (1) Heuristic predictor. We implement a simple heuristic-based model that considers all sentences containing the head or tail entity as supporting evidence. (2) Neural predictor. We also design a neural supporting evidence predictor. Given an entity pair and a predicted relation, sentences are first transformed into input representations by the concatenation of word embeddings and position embeddings, and then fed into a BiLSTM encoder for contextual representations. Inspired by Yang et al. (2018), we concatenate the output of the BiLSTM at the first and last positions with a trainable relation embedding to obtain a sentence’s representation, which is used to predict whether the sentence is adopted as supporting evidence for the given relation instance. As Table 7 shows, the neural predictor significantly outperforms heuristic-based baseline in predicting supporting evidence, which indicates the potential of RE models in joint relation and supporting evidence prediction.

Method	Dev	Test
Heuristic predictor	36.21	36.76
Neural predictor	<b>44.07</b>	<b>43.83</b>

Table 7: Performance of joint relation and supporting evidence prediction in F1 measurement (%).

**Discussion.** We can conclude from the above experimental results and analysis that document-level RE is more challenging than sentence-level RE and intensive efforts are needed to close the gap between the performance of RE models and human. We believe the following research directions are worth following: (1) Exploring models explicitly considering reasoning; (2) Designing more expressive model architectures for collecting and synthesizing inter-sentence information; (3) Leveraging distantly supervised data to improve the performance of document-level RE.



## 6 Related Work

A variety of datasets have been constructed for RE in recent years, which have greatly promoted the development of RE systems. [Hendrickx et al. \(2010\)](#), [Doddington et al. \(2004\)](#) and [Walker et al. \(2006\)](#) build human-annotated RE datasets with relatively limited relation types and instances. [Riedel et al. \(2010\)](#) automatically construct RE dataset by aligning plain text to KB via distant supervision, which suffers from wrong labeling problem. [Zhang et al. \(2017\)](#) and [Han et al. \(2018b\)](#) further combine external recommendations with human annotation to build large-scale high-quality datasets. **However, these RE datasets limit relations to single sentences.**

As documents provide richer information than sentences, moving research from sentence level to document level is a popular trend for many areas, including document-level event extraction ([Walker et al., 2006](#); [Mitamura et al., 2015, 2017](#)), fact extraction and verification ([Thorne et al., 2018](#)), reading comprehension ([Nguyen et al., 2016](#); [Joshi et al., 2017](#); [Lai et al., 2017](#)), sentiment classification ([Pang and Lee, 2004](#); [Prettenhofer and Stein, 2010](#)), summarization ([Nallapati et al., 2016](#)) and machine translation ([Zhang et al., 2018](#)). Recently, some document-level RE datasets have also been constructed. However, these datasets are either constructed via distant supervision ([Quirk and Poon, 2017](#); [Peng et al., 2017](#)) with inevitable wrong labeling problem, or limited in specific domain ([Li et al., 2016](#); [Peng et al., 2017](#)). In contrast, [DocRED](#) is constructed by crowd-workers with rich information, and is not limited in any specific domain, which makes it suitable to train and evaluate general-purpose document-level RE systems.

## 7 Conclusion

To promote RE systems from sentence level to document level, we present DocRED, a large-scale document-level RE dataset that features the data size, the requirement for reading and reasoning over multiple sentences, and the distantly supervised data offered for facilitating the development of weakly supervised document-level RE. Experiments show that human performance is significantly higher than RE baseline models, which suggests ample opportunity for future improvement.

## 8 Acknowledgement

This work is supported by the National Key Research and Development Program of China (No. 2018YFB1004503), the National Natural Science Foundation of China (NSFC No. 61572273) and China Association for Science and Technology (2016QNRC001). This work is also supported by the Pattern Recognition Center, WeChat AI, Tencent Inc. Han is also supported by 2018 Tencent Rhino-Bird Elite Training Program.

## References

- Rui Cai, Xiaodong Zhang, and Houfeng Wang. 2016. Bidirectional recurrent convolutional neural network for relation classification. In *Proceedings of ACL*, pages 756–765.
- Fenia Christopoulou, Makoto Miwa, and Sophia Ananiadou. 2018. A walk-based model on entity graphs for relation extraction. In *Proceedings of ACL*, pages 81–88.
- Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. 2019. BERT: pre-training of deep bidirectional transformers for language understanding. In *Proceedings of NAACL-HLT*.
- George Doddington, Alexis Mitchell, Mark Przybocki, Lance Ramshaw, Stephanie Strassel, and Ralph Weischedel. 2004. The automatic content extraction (ACE) program-tasks, data, and evaluation. In *Proceedings of LREC*, pages 837–840.
- Fredo Erxleben, Michael Günther, Markus Krötzsch, Julian Mendez, and Denny Vrandečić. 2014. Introducing wikidata to the linked data web. In *Proceedings of ISWC*, pages 50–65.
- Paolo Ferragina and Ugo Scaiella. 2010. TAGME: on-the-fly annotation of short text fragments (by wikipedia entities). In *Proceedings of CIKM*, pages 1625–1628.
- Xu Han, Pengfei Yu, Zhiyuan Liu, Maosong Sun, and Peng Li. 2018a. Hierarchical relation extraction with coarse-to-fine grained attention. In *Proceedings of EMNLP*, pages 2236–2245.
- Xu Han, Hao Zhu, Pengfei Yu, Ziyun Wang, Yuan Yao, Zhiyuan Liu, and Maosong Sun. 2018b. FewRel: A large-scale supervised few-shot relation classification dataset with state-of-the-art evaluation. In *Proceedings of EMNLP*, pages 4803–4809.
- Iris Hendrickx, Su Nam Kim, Zornitsa Kozareva, Preslav Nakov, Diarmuid Ó Séaghdha, Sebastian Padó, Marco Pennacchiotti, Lorenza Romano, and Stan Szpakowicz. 2010. SemEval-2010 task 8: Multi-way classification of semantic relations between pairs of nominals. In *Proceedings of ACL, the Workshop on Semantic Evaluation*, pages 33–38.

- Sepp Hochreiter and Jürgen Schmidhuber. 1997. [Long short-term memory](#). *Neural Computation*, 9:1735–1780.
- Raphael Hoffmann, Congle Zhang, Xiao Ling, Luke Zettlemoyer, and Daniel S. Weld. 2011. [Knowledge-based weak supervision for information extraction of overlapping relations](#). In *Proceedings of ACL*, pages 541–550.
- Mandar Joshi, Eunsol Choi, Daniel Weld, and Luke Zettlemoyer. 2017. [TriviaQA: A large scale distantly supervised challenge dataset for reading comprehension](#). In *Proceedings of ACL*, pages 1601–1611.
- Guokun Lai, Qizhe Xie, Hanxiao Liu, Yiming Yang, and Eduard Hovy. 2017. [RACE: Large-scale reading comprehension dataset from examinations](#). In *Proceedings of EMNLP*, pages 785–794.
- Omer Levy, Minjoon Seo, Eunsol Choi, and Luke Zettlemoyer. 2017. [Zero-shot relation extraction via reading comprehension](#). In *Proceedings of CoNLL*, pages 333–342.
- Jiao Li, Yueping Sun, Robin J. Johnson, Daniela Sciaky, Chih-Hsuan Wei, Robert Leaman, Allan Peter Davis, Carolyn J. Mattingly, Thomas C. Wiegers, and Zhiyong Lu. 2016. [BioCreative V CDR task corpus: a resource for chemical disease relation extraction](#). *Database*, pages 1–10.
- Yankai Lin, Shiqi Shen, Zhiyuan Liu, Huanbo Luan, and Maosong Sun. 2016. [Neural relation extraction with selective attention over instances](#). In *Proceedings of ACL*, pages 2124–2133.
- Mike Mintz, Steven Bills, Rion Snow, and Daniel Jurafsky. 2009. [Distant supervision for relation extraction without labeled data](#). In *Proceedings of ACL*, pages 1003–1011.
- Teruko Mitamura, Zhengzhong Liu, and Eduard H. Hovy. 2015. [Overview of TAC KBP 2015 event nugget track](#). In *Proceedings of TAC*, pages 16–17.
- Teruko Mitamura, Zhengzhong Liu, and Eduard H. Hovy. 2017. [Events detection, coreference and sequencing: What’s next? overview of the TAC KBP 2017 event track](#). In *Proceedings of TAC*.
- Ramesh Nallapati, Bowen Zhou, Cícero Nogueira dos Santos, Çağlar Gülçehre, and Bing Xiang. 2016. [Abstractive text summarization using sequence-to-sequence rnns and beyond](#). In *Proceedings of CoNLL*, pages 280–290.
- Tri Nguyen, Mir Rosenberg, Xia Song, Jianfeng Gao, Saurabh Tiwary, Rangan Majumder, and Li Deng. 2016. [MS MARCO: A human generated machine reading comprehension dataset](#). In *Proceedings of NIPS, the Workshop on Cognitive Computation*.
- Bo Pang and Lillian Lee. 2004. [A sentimental education: Sentiment analysis using subjectivity summarization based on minimum cuts](#). In *Proceedings of ACL*, pages 271–278.
- Nanyun Peng, Hoifung Poon, Chris Quirk, Kristina Toutanova, and Wen-tau Yih. 2017. [Cross-sentence n-ary relation extraction with graph LSTMs](#). *TACL*, 5:101–115.
- Jeffrey Pennington, Richard Socher, and Christopher Manning. 2014. [Glove: Global vectors for word representation](#). In *Proceedings of EMNLP*, pages 1532–1543.
- Peter Prettenhofer and Benno Stein. 2010. [Cross-language text classification using structural correspondence learning](#). In *Proceedings of ACL*, pages 1118–1127.
- Pengda Qin, Weiran Xu, and William Yang Wang. 2018. [Robust distant supervision relation extraction via deep reinforcement learning](#). In *Proceedings of ACL*, pages 2137–2147.
- Chris Quirk and Hoifung Poon. 2017. [Distant supervision for relation extraction beyond the sentence boundary](#). In *Proceedings of EACL*, pages 1171–1182.
- Sebastian Riedel, Limin Yao, and Andrew McCallum. 2010. [Modeling relations and their mentions without labeled text](#). In *Proceedings of ECML-PKDD*, pages 148–163.
- Cicero dos Santos, Bing Xiang, and Bowen Zhou. 2015. [Classifying relations by ranking with convolutional neural networks](#). In *Proceedings of ACL-IJCNLP*, pages 626–634.
- Richard Socher, Brody Huval, Christopher D. Manning, and Andrew Y. Ng. 2012. [Semantic compositionality through recursive matrix-vector spaces](#). In *Proceedings of EMNLP-CoNLL*, pages 1201–1211.
- Daniil Sorokin and Iryna Gurevych. 2017. [Context-aware representations for knowledge base relation extraction](#). In *Proceedings of EMNLP*, pages 1784–1789.
- Mihai Surdeanu, Julie Tibshirani, Ramesh Nallapati, and Christopher D. Manning. 2012. [Multi-instance multi-label learning for relation extraction](#). In *Proceedings of EMNLP-CoNLL*, pages 455–465.
- Kumutha Swampillai and Mark Stevenson. 2010. [Inter-sentential relations in information extraction corpora](#). In *Proceedings of LREC*, pages 2637–2641.
- James Thorne, Andreas Vlachos, Christos Christodoulopoulos, and Arpit Mittal. 2018. [FEVER: a large-scale dataset for fact extraction and verification](#). In *Proceedings of NAACL-HLT*, pages 809–819.

- Erik F Tjong Kim Sang and Fien De Meulder. 2003. [Introduction to the conll-2003 shared task: language-independent named entity recognition](#). In *Proceedings of NAACL-HLT*.
- Patrick Verga, Emma Strubell, and Andrew McCallum. 2018. [Simultaneously self-attending to all mentions for full-abstract biological relation extraction](#). In *Proceedings of NAACL-HLT*, pages 872–884.
- Denny Vrandečić and Markus Krötzsch. 2014. Wikidata: a free collaborative knowledgebase. *Communications of the ACM*.
- Christopher Walker, Stephanie Strassel, Julie Medero, and Kazuaki Maeda. 2006. ACE 2005 multilingual training corpus. *Linguistic Data Consortium, Philadelphia*, 57.
- Yi Wu, David Bamman, and Stuart Russell. 2017. [Adversarial training for relation extraction](#). In *Proceedings of EMNLP*, pages 1778–1783.
- Minguan Xiao and Cong Liu. 2016. [Semantic relation classification via hierarchical recurrent neural network with attention](#). In *Proceedings of COLING 2016*, pages 1254–1263.
- Zhilin Yang, Peng Qi, Saizheng Zhang, Yoshua Bengio, William W. Cohen, Ruslan Salakhutdinov, and Christopher D. Manning. 2018. [Hotpotqa: A dataset for diverse, explainable multi-hop question answering](#). In *Proceedings of EMNLP*, pages 2369–2380.
- Daojian Zeng, Kang Liu, Yubo Chen, and Jun Zhao. 2015. [Distant supervision for relation extraction via piecewise convolutional neural networks](#). In *Proceedings of EMNLP*, pages 1753–1762.
- Daojian Zeng, Kang Liu, Siwei Lai, Guangyou Zhou, and Jun Zhao. 2014. [Relation classification via convolutional deep neural network](#). In *Proceedings of COLING*, pages 2335–2344.
- Jiacheng Zhang, Huanbo Luan, Maosong Sun, Feifei Zhai, Jingfang Xu, Min Zhang, and Yang Liu. 2018. [Improving the transformer translation model with document-level context](#). In *Proceedings of EMNLP*, pages 533–542.
- Yuhao Zhang, Victor Zhong, Danqi Chen, Gabor Angeli, and Manning Christopher D. 2017. [Position-aware attention and supervised data improve slot filling](#). In *Proceedings of EMNLP*, pages 35–45.

## A Appendices

### A.1 Experimental Details

In this section, we provide more details of our experiments. To fairly compare the results of different models, we optimized all baselines using Adam, with learning rate of 0.001,  $\beta_1 = 0.9$ ,  $\beta_2 = 0.999$ . The other experimental hyper-parameters used in our experiments are shown in Table 8. Additionally, due to the document-level distance between entities, distances are first divided into several bins  $\{1, 2, \dots, 2^k\}$ , where each bin is associated with a trainable distance embedding.

Batch size	40
CNN hidden size	200
CNN window size	3
CNN dropout rate	0.5
LSTM hidden size	128
LSTM dropout rate	0.2
Word embedding dimension	100
Entity type embedding dimension	20
Coreference embedding dimension	20
Distance embedding dimension	20

Table 8: Hyper-parameter settings.

### A.2 Types of Named Entities

In this paper, we adapt the existing types of named entities used in Tjong Kim Sang and De Meulder (2003) to better serve DocRED. These types include “Person (PER)”, “Organization (ORG)”, “Location (LOC)”, “Time (TIME)”, “Number (NUM)”, and “other types (MISC)”. The types of named entities in DocRED and their covered contents are shown in Table 9.

Types	Content
PER	People, including fictional
ORG	Companies, universities, institutions, political or religious groups, etc.
LOC	Geographically defined locations, including mountains, waters, etc. Politically defined locations, including countries, cities, states, streets, etc. Facilities, including buildings, museums, stadiums, hospitals, factories, airports, etc.
TIME	Absolute or relative dates or periods.
NUM	Percents, money, quantities
MISC	Products, including vehicles, weapons, etc. Events, including elections, battles, sporting events, etc. Laws, cases, languages, etc

Table 9: Types of named entities in DocRED.

### A.3 List of Relations

We provide the list of relations in DocRED, including Wikidata IDs, relation names and descriptions from Wikidata in Table 10 and 11.



Wikidata ID	Name	Description
P6	head of government	head of the executive power of this town, city, municipality, state, country, or other governmental body
P17	country	sovereign state of this item; don't use on humans
P19	place of birth	most specific known (e.g. city instead of country, or hospital instead of city) birth location of a person, animal or fictional character
P20	place of death	most specific known (e.g. city instead of country, or hospital instead of city) death location of a person, animal or fictional character
P22	father	male parent of the subject. For stepfather, use "stepparent" (P3448)
P25	mother	female parent of the subject. For stepmother, use "stepparent" (P3448)
P26	spouse	the subject has the object as their spouse (husband, wife, partner, etc.). Use "unmarried partner" (P451) for non-married companions
P27	country of citizenship	the object is a country that recognizes the subject as its citizen
P30	continent	continent of which the subject is a part
P31	instance of	that class of which this subject is a particular example and member. (Subject typically an individual member with Proper Name label.) Different from P279 (subclass of)
P35	head of state	official with the highest formal authority in a country/state
P36	capital	primary city of a country, state or other type of administrative territorial entity
P37	official language	language designated as official by this item
P39	position held	subject currently or formerly holds the object position or public office
P40	child	subject has the object in their family as their offspring son or daughter (independently of their age)
P50	author	main creator(s) of a written work (use on works, not humans)
P54	member of sports team	sports teams or clubs that the subject currently represents or formerly represented
P57	director	director(s) of this motion picture, TV-series, stageplay, video game or similar
P58	screenwriter	author(s) of the screenplay or script for this work
P69	educated at	educational institution attended by the subject
P86	composer	person(s) who wrote the music; also use P676 for lyricist
P102	member of political party	the political party of which this politician is or has been a member
P108	employer	person or organization for which the subject works or worked
P112	founded by	founder or co-founder of this organization, religion or place
P118	league	league in which team or player plays or has played in
P123	publisher	organization or person responsible for publishing books, periodicals, games or software
P127	owned by	owner of the subject
P131	located in the administrative territorial entity	the item is located on the territory of the following administrative entity. Use P276 (location) for specifying the location of non-administrative places and for items about events
P136	genre	a creative work's genre or an artist's field of work (P101). Use main subject (P921) to relate creative works to their topic
P137	operator	person or organization that operates the equipment, facility, or service; use country for diplomatic missions
P140	religion	religion of a person, organization or religious building, or associated with this subject
P150	contains administrative territorial entity	(list of) direct subdivisions of an administrative territorial entity
P155	follows	immediately prior item in some series of which the subject is part. Use P1365 (replaces) if the preceding item was replaced, e.g. political offices, states and there is no identity between precedent and following geographic unit
P156	followed by	immediately following item in some series of which the subject is part. Use P1366 (replaced by) if the item is replaced, e.g. political offices, states
P159	headquarters location	specific location where an organization's headquarters is or has been situated
P161	cast member	actor performing live for a camera or audience [use "character role" (P453) and/or "name of the character role" (P4633) as qualifiers] [use "voice actor" (P725) for voice-only role]
P162	producer	producer(s) of this film or music work (film: not executive producers, associate producers, etc.) [use P272 to refer to the production company]
P166	award received	award or recognition received by a person, organisation or creative work
P170	creator	maker of a creative work or other object (where no more specific property exists)
P171	parent taxon	closest parent taxon of the taxon in question
P172	ethnic group	subject's ethnicity (consensus is that a VERY high standard of proof is needed for this field to be used. In general this means 1) the subject claims it him/herself, or 2) it is widely agreed on by scholars, or 3) is fictional and portrayed as such).
P175	performer	performer involved in the performance or the recording of a work
P176	manufacturer	manufacturer or producer of this product
P178	developer	organisation or person that developed this item
P179	series	subject is part of a series, whose sum constitutes the object
P190	sister city	twin towns, sister cities, twinned municipalities and other localities that have a partnership or cooperative agreement, either legally or informally acknowledged by their governments
P194	legislative body	legislative body governing this entity; political institution with elected representatives, such as a parliament/legislature or council
P205	basin country	country that have drainage to/from or border the body of water

Table 10: Relation list (I), including Wikidata IDs, names and descriptions of relations in DocRED.

Wikidata ID	Name	Description
P206	located in or next to body of water	sea, lake or river
P241	military branch	branch to which this military unit, award, office, or person belongs, e.g. Royal Navy
P264	record label	brand and trademark associated with the marketing of subject music recordings and music videos
P272	production company	company that produced this film, audio or performing arts work
P276	location	location of the item, physical object or event is within. In case of an administrative entity use P131. In case of a distinct terrain feature use P706.
P279	subclass of	all instances of these items are instances of those items; this item is a class (subset) of that item. Not to be confused with P31 (instance of)
P355	subsidiary	subsidiary of a company or organization, opposite of parent company (P749)
P361	part of	object of which the subject is a part. Inverse property of "has part" (P527). See also "has parts of the class" (P2670).
P364	original language of work	language in which a film or a performance work was originally created. Deprecated for written works; use P407 ("language of work or name") instead.
P400	platform	platform for which a work has been developed or released / specific platform version of a software developed
P403	mouth of the watercourse	the body of water to which the watercourse drains
P449	original network	network(s) the radio or television show was originally aired on, including
P463	member of	organization or club to which the subject belongs. Do not use for membership in ethnic or social groups, nor for holding a position such as a member of parliament (use P39 for that).
P488	chairperson	presiding member of an organization, group or body
P495	country of origin	country of origin of the creative work or subject item
P527	has part	part of this subject. Inverse property of "part of" (P361).
P551	residence	the place where the person is, or has been, resident
P569	date of birth	date on which the subject was born
P570	date of death	date on which the subject died
P571	inception	date or point in time when the organization/subject was founded/created
P576	dissolved, abolished or demolished	date or point in time on which an organisation was dissolved/disappeared or a building demolished; see also discontinued date (P2669)
P577	publication date	date or point in time a work is first published or released
P580	start time	indicates the time an item begins to exist or a statement starts being valid
P582	end time	indicates the time an item ceases to exist or a statement stops being valid
P585	point in time	time and date something took place, existed or a statement was true
P607	conflict	battles, wars or other military engagements in which the person or item participated
P674	characters	characters which appear in this item (like plays, operas, operettas, books, comics, films, TV series, video games)
P676	lyrics by	author of song lyrics; also use P86 for music composer
P706	located on terrain feature	located on the specified landform. Should not be used when the value is only political/administrative (provinces, states, countries, etc.). Use P131 for administrative entity.
P710	participant	person, group of people or organization (object) that actively takes/took part in the event (subject). Preferably qualify with "object has role" (P3831). Use P1923 for team participants.
P737	influenced by	this person, idea, etc. is informed by that other person, idea, etc., e.g. "Heidegger was influenced by Aristotle".
P740	location of formation	location where a group or organization was formed
P749	parent organization	parent organization of an organisation, opposite of subsidiaries (P355)
P800	notable work	notable scientific, artistic or literary work, or other work of significance among subject's works
P807	separated from	subject was founded or started by separating from identified object
P840	narrative location	the narrative of the work is set in this location
P937	work location	location where persons were active
P1001	applies to jurisdiction	the item (an institution, law, public office ...) belongs to or has power over or applies to the value (a territorial jurisdiction: a country, state, municipality, ...)
P1056	product or material produced	material or product produced by a government agency, business, industry, facility, or process
P1198	unemployment rate	portion of a workforce population that is not employed
P1336	territory claimed by	administrative divisions that claim control of a given area
P1344	participant of	event a person or an organization was a participant in, inverse of P710 or P1923
P1365	replaces	person or item replaced. Use P1398 (structure replaces) for structures. Use P155 (follows) if the previous item was not replaced or if predecessor and successor are identical.
P1366	replaced by	person or item which replaces another. Use P156 (followed by) if the item is not replaced (e.g. books in a series), nor identical
P1376	capital of	country, state, department, canton or other administrative division of which the municipality is the governmental seat
P1412	languages spoken, written or signed	language(s) that a person speaks or writes, including the native language(s)
P1441	present in work	work in which this fictional entity (Q14897293) or historical person is present
P3373	sibling	the subject has the object as their sibling (brother, sister, etc.). Use "relative" (P1038) for siblings-in-law (brother-in-law, sister-in-law, etc.) and step-siblings (step-brothers, step-sisters, etc.)

Table 11: Relation list (II), including Wikidata IDs, names and descriptions of relations in DocRED.